

SURBHI GOEL

<https://www.surbhigoel.com>

[first name][last initial]@cis.upenn.edu

EDUCATION

The University of Texas at Austin

August 2015 - June 2020

M.S. and Ph.D. in Computer Science

Advisor: Adam R. Klivans

Committee: Alex Dimakis, Raghu Meka, Eric Price

Dissertation: [Towards Provably Efficient Algorithms for Learning Neural Networks](#)

Received the Bert Kay dissertation award

Indian Institute of Technology, Delhi

July 2011 - May 2015

B.Tech. in Computer Science and Engineering

APPOINTMENTS

University of Pennsylvania, Philadelphia, PA

January 2023 - Present

Magerman Term Assistant Professor, Computer and Information Science

Microsoft Research, New York, NY

July 2020 - December 2022

Postdoctoral Researcher, Machine Learning Group

Institute for Advanced Study, Princeton, NJ

January - May 2020

Visiting Graduate Student, Theoretical Machine Learning Program

Simons Institute for Theory of Computing, Berkeley, CA

May - August 2019

Research Fellow, Foundations of Deep Learning Program

RESEARCH INTERESTS

My research is on the theoretical aspects of the modern practice of machine learning, where my goal is to develop the next generation of principled machine learning methods. In the pursuit of this goal, my work focuses on quantifying the computational and statistical aspects of state-of-the-art deep learning methods, and expanding the toolbox of current algorithms using new theoretically grounded insights.

AWARDS AND FELLOWSHIPS

- 2020 Bert Kay Dissertation Award for best dissertation in CS at UT Austin
- 2019 Rising Stars in ML by University of Maryland
- 2019 Rising Stars in EECS by UIUC
- 2019 The University of Texas at Austin Graduate Dean's Prestigious Fellowship Supplement
- 2019 J.P. Morgan AI PhD Fellowship
- 2019 Simons-Berkeley Research Fellowship for Foundations of Deep Learning program
- 2018 The University of Texas at Austin Graduate Continuing Bruton Fellowship
- 2017 The University of Texas at Austin Graduate School Summer Fellowship
- 2015 ICIM Stay Ahead Award for Undergraduate Thesis
- 2015 Suresh Chandra Memorial Trust Award for Undergraduate Thesis
- 2011 Aditya Birla Scholarship awarded to 12 students from all over India
- 2011 OPJEM Scholarship awarded to 1 out of 850 students in the batch at IIT Delhi
- 2011 All India Rank 37 (Rank 2 in girls) in IITJEE among 450,000 students
- 2010 National Mathematics Olympiad finalist (1 out of 30 from all over India)

PUBLICATIONS

* indicates α - β (alphabetical) ordering.

WORKING PAPERS

Kan Xu, Hamsa Bastani, **Surbhi Goel**, Osbert Bastani
Stochastic Bandits with ReLU Neural Networks
In submission, 2023

CONFERENCE PAPERS

Surbhi Goel*, Steve Hanneke*, Shay Moran*, Abhishek Shetty*
Adversarial Resilience in Sequential Prediction via Abstention
Neural Information Processing Systems (NeurIPS) 2023

Benjamin L. Edelman*, **Surbhi Goel***, Sham M. Kakade*, Eran Malach*, Cyril Zhang*
Pareto Frontiers in Neural Feature Learning: Data, Compute, Width, and Luck
Neural Information Processing Systems (NeurIPS) 2023
Selected as a spotlight presentation

Bingbin Liu, Jordan T. Ash, **Surbhi Goel**, Akshay Krishnamurthy, Cyril Zhang
Exposing Attention Glitches with Flip-Flop Language Modeling
Neural Information Processing Systems (NeurIPS) 2023
Selected as a spotlight presentation

Sitan Chen*, Zehao Dou*, **Surbhi Goel***, Adam R. Klivans*, Raghu Meka*
Learning Narrow One-Hidden-Layer ReLU Networks
Conference on Learning Theory (COLT) 2023

Bingbin Liu, Jordan T. Ash, **Surbhi Goel**, Akshay Krishnamurthy, Cyril Zhang
Transformers Learn Shortcuts to Automata
International Conference on Learning Representations (ICLR) 2023
Selected as a notable top-5% paper

Surbhi Goel*, Sham M. Kakade*, Adam T. Kalai*, Cyril Zhang*
Recurrent Convolutional Neural Networks Learn Succinct Learning Algorithms
Neural Information Processing Systems (NeurIPS) 2022

Boaz Barak*, Benjamin L. Edelman*, **Surbhi Goel***, Sham M. Kakade*, Eran Malach*, Cyril Zhang*
Hidden Progress in Deep Learning: SGD Learns Parities Near the Computational Limit
Neural Information Processing Systems (NeurIPS) 2022

Benjamin L. Edelman*, **Surbhi Goel***, Sham M. Kakade*, Cyril Zhang*
Inductive Biases and Variable Creation in Self-Attention Mechanisms
International Conference on Machine Learning (ICML) 2022

Nikunj Saunshi, Jordan T. Ash, **Surbhi Goel**, Dipendra Misra, Cyril Zhang, Sanjeev Arora, Sham M. Kakade, Akshay Krishnamurthy
Understanding Contrastive Learning Requires Incorporating Inductive Biases
International Conference on Machine Learning (ICML) 2022

Jordan T. Ash, Cyril Zhang, **Surbhi Goel**, Akshay Krishnamurthy, Sham M. Kakade
Anti-Concentrated Confidence Bonuses For Scalable Exploration
International Conference on Learning Representations (ICLR) 2022

Jordan T. Ash*, **Surbhi Goel***, Akshay Krishnamurthy*, Dipendra Misra*
Investigating the Role of Negatives in Contrastive Representation Learning
International Conference on Artificial Intelligence and Statistics (AISTATS) 2022

Jordan T. Ash, **Surbhi Goel**, Akshay Krishnamurthy, Sham M. Kakade
Gone Fishing: Neural Active Learning with Fisher Embeddings
Neural Information Processing Systems (NeurIPS) 2021

Naman Agarwal*, **Surbhi Goel***, Cyril Zhang*
Acceleration via Fractal Learning Rate Schedules
International Conference on Machine Learning (ICML) 2021

Anthimos-Vardis Kandiros, Yuval Dagan, Nishanth Dikkala, **Surbhi Goel**, Constantinos Daskalakis
Statistical Estimation from Dependent Data
International Conference on Machine Learning (ICML) 2021

Surbhi Goel*, Adam R. Klivans*, Pasin Manurangsi*, Daniel Reichman*
Tight Hardness Results for Learning One-Layer ReLU Networks
Innovations in Theoretical Computer Science (ITCS) 2021

Surbhi Goel*, Adam R. Klivans*, Frederic Koehler*
From Boltzmann Machines to Neural Networks and Back Again
Neural Information Processing Systems (NeurIPS) 2020

Surbhi Goel*, Aravind Gollakota*, Adam R., Klivans*
Statistical-Query Lower Bounds via Functional Gradients
Neural Information Processing Systems (NeurIPS) 2020

Surbhi Goel*, Aravind Gollakota*, Zhihan Jin*, Sushrut Karmalkar*, Adam R. Klivans*
Superpolynomial Lower Bounds for Learning One-Layer Neural Networks using Gradient Descent
International Conference on Machine Learning (ICML) 2020

Omar Montasser, **Surbhi Goel**, Ilias Diakonikolas, Nathan Srebro
Efficiently Learning Adversarially Robust Halfspaces with Noise
International Conference on Machine Learning (ICML) 2020

Jessica Hoffmann, Soumya Basu, **Surbhi Goel**, Constantine Caramanis
Learning Mixtures of Graphs from Epidemic Cascades
International Conference on Machine Learning (ICML) 2020

Ilias Diakonikolas*, **Surbhi Goel***, Sushrut Karmalkar*, Adam R. Klivans*, Mahdi Soltanolkotabi*
Approximation Schemes for ReLU Regression
Conference on Learning Theory (COLT) 2020

Surbhi Goel
Learning Ising and Potts Models with Latent Variables
International Conference on Artificial Intelligence and Statistics (AISTATS) 2020

Surbhi Goel*, Sushrut Karmalkar*, Adam R. Klivans*

Time/Accuracy Trade-offs for Learning a ReLU with respect to Gaussian Marginals

Neural Information Processing Systems (NeurIPS) 2019

Selected for a spotlight presentation

Surbhi Goel*, Daniel Kane*, Adam R. Klivans*

Learning Ising Models with Independent Failures

Conference on Learning Theory (COLT) 2019

Surbhi Goel*, Adam R. Klivans*

Learning Neural Networks with Two Nonlinear Layers in Polynomial Time

Conference on Learning Theory (COLT) 2019

Surbhi Goel*, Adam R. Klivans*, Raghu Meka*

Learning One Convolutional Layer with Overlapping Patches

International Conference on Machine Learning (ICML) 2018

Selected for a full oral presentation

Surbhi Goel*, Adam R. Klivans*

Eigenvalue Decay Implies Polynomial-Time Learnability for Neural Networks

Neural Information Processing Systems (NeurIPS) 2017

Surbhi Goel*, Varun Kanade*, Adam R. Klivans*, Justin Thaler*

Reliably Learning ReLU in Polynomial Time

Conference on Learning Theory (COLT) 2017

WORKSHOP PAPERS

Bingbin Liu, Jordan T. Ash, **Surbhi Goel**, Akshay Krishnamurthy, Cyril Zhang

Exposing Attention Glitches with Flip-Flop Language Modeling

Challenges of Deploying Generative AI Workshop, International Conference on Machine Learning (ICML) 2023

Knowledge and Logical Reasoning in the Era of Data-driven Learning Workshop, International Conference on Machine Learning (ICML) 2023

Jessica Hoffmann, Soumya Basu, **Surbhi Goel**, Constantine Caramanis

Disentangling Mixtures of Epidemics on Graphs

Graph Representation Learning, Neural Information Processing Systems (NeurIPS) 2019

Surbhi Goel*, Adam R. Klivans*

Learning Depth-Three Neural Networks in Polynomial Time

Deep Learning: Bridging Theory and Practice, Neural Information Processing Systems (NeurIPS) 2017

Surbhi Goel*, Varun Kanade*, Adam R. Klivans*, Justin Thaler*

Reliably Learning ReLU in Polynomial Time

Optimization for Machine Learning (OPT), Neural Information Processing Systems (NeurIPS) 2016

Selected for an oral presentation

UNPUBLISHED MANUSCRIPTS

Surbhi Goel*, Rina Panigrahy*

Learning Two layer Networks with Multinomial Activation and High Thresholds

Manuscript, 2019

Matthew Jordan, Naren Manoj, **Surbhi Goel**, Alexandros Dimakis
Quantifying Perceptual Distortion of Adversarial Examples
Manuscript, 2019

Simon Du*, **Surbhi Goel***
Improved Learning of One-hidden-layer Convolutional Neural Networks with Overlaps
Manuscript, 2018.

INVITED TALKS

Beyond Worst-case Sequential Prediction: Adversarial Robustness via Abstention

Math Machine Learning seminar at MPI MIS + UCLA

August 2023

FODSI Workshop on Computational Complexity of Statistical Problems at MIT

June 2023

Thinking fast with Transformers - Algorithmic Reasoning via Shortcuts

Youth in High Dimensions, Trieste, Italy

May 2023

MaD Seminar at NYU

April 2023

ASSET Seminar at UPenn

April 2023

Sparse Feature Emergence in Deep Learning

Symposium on New Directions in Theoretical Machine Learning [[slides](#)]

September 2022

What Functions do Self-attention Blocks Prefer to Represent?

Demystifying Attention-based Architectures in Deep Learning

Joint IFML/Data-Driven Decision Processes Workshop at Simons Institute

October 2022

ML Foundations Seminar at MSR Redmond

August 2022

Workshop on Algorithms for Learning and Economics (WALE) in Greece

June 2022

ML Symposium at USC

December 2021

ELLIS Talk Series at IST Austria

December 2021

Learning Theory Workshop at Google

October 2021

The Hidden Progress Behind Loss Curves

Workshop on Learning: Optimization and Stochastics at EPFL

July 2022

Principled Algorithm Design in the Era of Deep Learning

CS/CSE Colloquium at NYU Courant/Tandon

April 2022

CS Colloquium at UW-Madison

March 2022

CS Colloquium at Halicioglu Data Science Institute

March 2022

CS Colloquium at UMD

February 2022

SCS Talk at CMU

February 2022

CS Colloquium at Duke

February 2022

CIS Colloquium at UPenn

February 2022

CS Colloquium at Cornell

February 2022

Talks at TTIC

February 2022

Computational Barriers For Learning Some Generalized Linear Models

Information-Computation Trade-offs Workshop at Simons Institute [[video](#)][[slides](#)]

September 2021

Computational Complexity of ReLU Regression

The Multifaceted Complexity of Machine Learning Workshop at IMSI [[video](#)]

April 2021

Computational Complexity of Learning Neural Networks over Gaussian Marginals

<i>MIC Seminar at NYU</i>	<i>May 2020</i>
<i>Algorithms Seminar at Duke University</i>	<i>October 2020</i>
<i>ML Theory Seminar at Harvard University [video]</i>	<i>October 2020</i>
<i>ARC Colloquium at Georgia Tech</i>	<i>November 2020</i>
<i>IDEAL Seminar at TTIC</i>	<i>November 2020</i>
<i>TOC Colloquium at MIT</i>	<i>December 2020</i>
<i>SILO Seminar at UW-Madison</i>	<i>January 2020</i>
<i>Statistics Seminar at Stanford University</i>	<i>July 2021</i>

Approximation Schemes for ReLU Regression

<i>Deep Learning Program Reunion at Simons Institute</i>	<i>August 2020</i>
--	--------------------

Provably Efficient Algorithms for Learning Neural Networks

<i>Microsoft Research New York</i>	<i>February 2020</i>
<i>Microsoft Research New England</i>	<i>February 2020</i>
<i>Microsoft Research Redmond</i>	<i>February 2020</i>

Time/Accuracy Tradeoffs for Learning a ReLU wrt Gaussian Marginals

<i>Spotlight Talk at Neural Information Processing Systems (NeurIPS)</i>	<i>December 2019</i>
--	----------------------

Exploring Surrogate Losses for Learning Neural Networks

<i>TTIC Young Researcher Seminar Series</i>	<i>December 2019</i>
---	----------------------

Efficiently Learning Simple Neural Networks

<i>Rising Star in ML Talk at University of Maryland</i>	<i>September 2019</i>
---	-----------------------

Learning Ising Models with Independent Failures

<i>Research Fellows Talk at Simons Institute</i>	<i>July 2019</i>
--	------------------

Efficiently Learning Simple Convolutional Networks

<i>China Theory Week at Tsinghua University</i>	<i>September 2019</i>
---	-----------------------

Learning One Convolutional Layer with Overlapping Patches

<i>Google Research Theory Reading Group</i>	<i>June 2018</i>
---	------------------

Reliably Learning the ReLU in Polynomial Time

<i>OPT-ML Workshop at Neural Information Processing Systems (NeurIPS)</i>	<i>December 2016</i>
---	----------------------

WORK EXPERIENCE

Google, Mountain View CA

Research Intern

May - August 2018

Supervisor: Rina Panigrahy

Dell, Round Rock TX

Research Intern

June - August 2017

Google, New York, NY

Research Intern

May - August 2016

Supervisor: Natalia Ponomareva

Google, Mountain View CA

Software Engineering Intern

May - August 2014

Supervisor: Neha Jha

OUTREACH

Co-founder

2020-Present

Learning Theory Alliance (LeT-All)

Co-organizing the *Fall 2023 Mentoring Workshop*

Co-organized the *Fall 2022 Mentoring Workshop* in collaboration with FODSI

Co-organized the *COLT 2022 Mentoring Panel*

Co-organized the *ALT 2022 Mentoring Workshop*

Co-organized the *Graduate Applications Support Program* in collaboration with WiML-T

Co-organized the *COLT 2021 Mentoring Workshop*

Co-organized the *ALT 2021 Mentoring Workshop*

Mentor

Women in Machine Learning Theory (WiML-T) Mentoring Program

2021-Present

UT Austin's Women in CS (GWC-WiCS) Mentoring Program

2018-19

Panelist

WiML Un-Workshop, ICML 2022

July 2022

New Horizons in Theoretical Computer Science

June 2022

VMware Nirman for Women in Tech

January 2021

SERVICE ROLES

Office Hours Chair

2023-24

International Conference on Learning Representations (ICLR) 2024

Workshop Co-organizer

2023

Mathematics of Modern Machine Learning (M3L) at NeurIPS 2023

Virtual Experience Chair

2023

Conference on Learning Theory (COLT) 2023

Online Experience Chair

2021

Conference on Learning Theory (COLT) 2021

Co-organized the virtual part of the hybrid conference, including the 2-day virtual-only program

Seminar Co-organizer

2020-21

One World Machine Learning Seminar Series

Treasurer

2016-17

Graduate Representative Association of Computer Sciences (GRACS)

Program Committee

International Conference on Algorithmic Learning Theory (ALT)

2021/22/23

Conference on Learning Theory (COLT)

2021/22

International Conference on Artificial Intelligence and Statistics (AISTATS) (area chair)

2023

Neural Information Processing Systems (NeurIPS) (area chair)

2023

International Conference on Algorithmic Learning Theory (ALT) (senior program committee)

2024

Conference Reviewing

<i>Symposium on Theory of Computing (STOC)</i>	2019/20/21
<i>Neural Information Processing Systems (NeurIPS)</i>	2018 (top 30%)/20/21
<i>Conference on Learning Theory (COLT)</i>	2018/19/20
<i>International Conference on Learning Representations (ICLR)</i>	2019/20/23
<i>Symposium on Discrete Algorithms (SODA)</i>	2020/23
<i>Foundations of Computer Science (FOCS)</i>	2020/22
<i>International Conference on Machine Learning (ICML)</i>	2019 (top 5%)

Journal Reviewing

<i>Journal of Machine Learning Research</i>	2021/22
<i>IEEE Transactions on Information Theory</i>	2020