

SURBHI GOEL

<https://www.surbhigoel.com>

[first name][last initial]@cis.upenn.edu

EDUCATION

The University of Texas at Austin

August 2015 - June 2020

M.S. and Ph.D. in Computer Science

Advisor: Adam R. Klivans

Committee: Alex Dimakis, Raghu Meka, Eric Price

Dissertation: [Towards Provably Efficient Algorithms for Learning Neural Networks](#)

Received the Bert Kay dissertation award

Indian Institute of Technology, Delhi

July 2011 - May 2015

B.Tech. in Computer Science and Engineering

APPOINTMENTS

University of Pennsylvania, Philadelphia, PA

January 2023 - Present

Magerman Term Assistant Professor, Computer and Information Science

Microsoft Research, New York, NY

July 2020 - December 2022

Postdoctoral Researcher, Machine Learning Group

Institute for Advanced Study, Princeton, NJ

January - May 2020

Visiting Graduate Student, Theoretical Machine Learning Program

Simons Institute for Theory of Computing, Berkeley, CA

May - August 2019

Research Fellow, Foundations of Deep Learning Program

RESEARCH INTERESTS

My research is on the theoretical aspects of the modern practice of machine learning, where my goal is to develop the next generation of principled machine learning methods. In the pursuit of this goal, my work focuses on quantifying the computational and statistical aspects of state-of-the-art deep learning methods, and expanding the toolbox of current algorithms using new theoretically grounded insights.

AWARDS AND FELLOWSHIPS

- 2024 OpenAI Superalignment Fast Grant (\$150,000)
- 2023 Microsoft Accelerate Foundation Models Research Award (\$25,000)
- 2020 Bert Kay Dissertation Award for best dissertation in CS at UT Austin
- 2019 Rising Star in ML by University of Maryland
- 2019 Rising Star in EECS by UIUC
- 2019 J.P. Morgan AI PhD Fellowship
- 2019 Simons-Berkeley Research Fellowship for Foundations of Deep Learning program
- 2018 The University of Texas at Austin Graduate Continuing Bruton Fellowship
- 2017 The University of Texas at Austin Graduate School Summer Fellowship
- 2015 ICIM Stay Ahead Award and Suresh Chandra Memorial Trust Award for Undergraduate Thesis
- 2011 Aditya Birla Scholarship
- 2011 OPJEM Scholarship
- 2011 All India Rank 37 (*Rank 2 amongst women applicants*) in IITJEE among 450,000 students
- 2010 Indian National Mathematics Olympiad Top 30

PUBLICATIONS

* indicates α - β (alphabetical) ordering.

PREPRINTS

Mahdi Sabbaghi, George J. Pappas, Hamed Hassani, **Surbhi Goel**

Encoding Structural Symmetry is Key for Length Generalization in Arithmetic Tasks

In submission, 2024

Anton Xue, Avishree Khare, Rajeev Alur, **Surbhi Goel**, Eric Wong

Logicbreaks: A Framework for Understanding Subversion of Rule-based Inference

In submission, 2024

CONFERENCE PAPERS

Surbhi Goel*, Abhishek Shetty*, Konstantinos Stavropoulos*, Arsen Vasilyan*

Tolerant Algorithms for Learning with Arbitrary Covariate Shift

Neural Information Processing Systems (NeurIPS), 2024

Selected as a spotlight presentation

Benjamin L. Edelman*, Ezra Edelman*, **Surbhi Goel***, Eran Malach*, Nikolaos Tsilivis*

The Evolution of Statistical Induction Heads: In-Context Learning Markov Chains

Neural Information Processing Systems (NeurIPS), 2024

GuanWen Qiu, Da Kuang, **Surbhi Goel**

Complexity Matters: Feature Learning in the Presence of Spurious Correlations

International Conference on Machine Learning (ICML) 2024

Kan Xu, Hamsa Bastani, **Surbhi Goel**, Osbert Bastani

Stochastic Bandits with ReLU Neural Networks

International Conference on Machine Learning (ICML) 2024

Surbhi Goel*, Steve Hanneke*, Shay Moran*, Abhishek Shetty*

Adversarial Resilience in Sequential Prediction via Abstention

Neural Information Processing Systems (NeurIPS) 2023

Benjamin L. Edelman*, **Surbhi Goel***, Sham M. Kakade*, Eran Malach*, Cyril Zhang*

Pareto Frontiers in Neural Feature Learning: Data, Compute, Width, and Luck

Neural Information Processing Systems (NeurIPS) 2023

Selected as a spotlight presentation

Bingbin Liu, Jordan T. Ash, **Surbhi Goel**, Akshay Krishnamurthy, Cyril Zhang

Exposing Attention Glitches with Flip-Flop Language Modeling

Neural Information Processing Systems (NeurIPS) 2023

Selected as a spotlight presentation

Sitan Chen*, Zehao Dou*, **Surbhi Goel***, Adam R. Klivans*, Raghu Meka*

Learning Narrow One-Hidden-Layer ReLU Networks

Conference on Learning Theory (COLT) 2023

Bingbin Liu, Jordan T. Ash, **Surbhi Goel**, Akshay Krishnamurthy, Cyril Zhang

Transformers Learn Shortcuts to Automata

International Conference on Learning Representations (ICLR) 2023

Selected as a notable top-5% paper

Surbhi Goel*, Sham M. Kakade*, Adam T. Kalai*, Cyril Zhang*

Recurrent Convolutional Neural Networks Learn Succinct Learning Algorithms

Neural Information Processing Systems (NeurIPS) 2022

Boaz Barak*, Benjamin L. Edelman*, **Surbhi Goel***, Sham M. Kakade*, Eran Malach*, Cyril Zhang*

Hidden Progress in Deep Learning: SGD Learns Parities Near the Computational Limit

Neural Information Processing Systems (NeurIPS) 2022

Benjamin L. Edelman*, **Surbhi Goel***, Sham M. Kakade*, Cyril Zhang *

Inductive Biases and Variable Creation in Self-Attention Mechanisms

International Conference on Machine Learning (ICML) 2022

Nikunj Saunshi, Jordan T. Ash, **Surbhi Goel**, Dipendra Misra, Cyril Zhang, Sanjeev Arora, Sham M. Kakade, Akshay Krishnamurthy

Understanding Contrastive Learning Requires Incorporating Inductive Biases

International Conference on Machine Learning (ICML) 2022

Jordan T. Ash, Cyril Zhang, **Surbhi Goel**, Akshay Krishnamurthy, Sham M. Kakade

Anti-Concentrated Confidence Bonuses For Scalable Exploration

International Conference on Learning Representations (ICLR) 2022

Jordan T. Ash*, **Surbhi Goel***, Akshay Krishnamurthy*, Dipendra Misra*

Investigating the Role of Negatives in Contrastive Representation Learning

International Conference on Artificial Intelligence and Statistics (AISTATS) 2022

Jordan T. Ash, **Surbhi Goel**, Akshay Krishnamurthy, Sham M. Kakade

Gone Fishing: Neural Active Learning with Fisher Embeddings

Neural Information Processing Systems (NeurIPS) 2021

Naman Agarwal*, **Surbhi Goel***, Cyril Zhang*

Acceleration via Fractal Learning Rate Schedules

International Conference on Machine Learning (ICML) 2021

Anthimos-Vardis Kandiros, Yuval Dagan, Nishanth Dikkala, **Surbhi Goel**, Constantinos Daskalakis

Statistical Estimation from Dependent Data

International Conference on Machine Learning (ICML) 2021

Surbhi Goel*, Adam R. Klivans*, Pasin Manurangsi*, Daniel Reichman*

Tight Hardness Results for Learning One-Layer ReLU Networks

Innovations in Theoretical Computer Science (ITCS) 2021

Surbhi Goel*, Adam R. Klivans*, Frederic Koehler*

From Boltzmann Machines to Neural Networks and Back Again

Neural Information Processing Systems (NeurIPS) 2020

Surbhi Goel*, Aravind Gollakota*, Adam R., Klivans*

Statistical-Query Lower Bounds via Functional Gradients

Neural Information Processing Systems (NeurIPS) 2020

Surbhi Goel*, Aravind Gollakota*, Zhihan Jin*, Sushrut Karmalkar*, Adam R. Klivans*

Superpolynomial Lower Bounds for Learning One-Layer Neural Networks using Gradient Descent

International Conference on Machine Learning (ICML) 2020

Omar Montasser, **Surbhi Goel**, Ilias Diakonikolas, Nathan Srebro

Efficiently Learning Adversarially Robust Halfspaces with Noise

International Conference on Machine Learning (ICML) 2020

Jessica Hoffmann, Soumya Basu, **Surbhi Goel**, Constantine Caramanis

Learning Mixtures of Graphs from Epidemic Cascades

International Conference on Machine Learning (ICML) 2020

Ilias Diakonikolas*, **Surbhi Goel***, Sushrut Karmalkar*, Adam R. Klivans*, Mahdi Soltanolkotabi*

Approximation Schemes for ReLU Regression

Conference on Learning Theory (COLT) 2020

Surbhi Goel

Learning Ising and Potts Models with Latent Variables

International Conference on Artificial Intelligence and Statistics (AISTATS) 2020

Surbhi Goel*, Sushrut Karmalkar*, Adam R. Klivans*

Time/Accuracy Trade-offs for Learning a ReLU with respect to Gaussian Marginals

Neural Information Processing Systems (NeurIPS) 2019

Selected for a spotlight presentation

Surbhi Goel*, Daniel Kane*, Adam R. Klivans*

Learning Ising Models with Independent Failures

Conference on Learning Theory (COLT) 2019

Surbhi Goel*, Adam R. Klivans*

Learning Neural Networks with Two Nonlinear Layers in Polynomial Time

Conference on Learning Theory (COLT) 2019

Surbhi Goel*, Adam R. Klivans*, Raghu Meka*

Learning One Convolutional Layer with Overlapping Patches

International Conference on Machine Learning (ICML) 2018

Selected for a full oral presentation

Surbhi Goel*, Adam R. Klivans*

Eigenvalue Decay Implies Polynomial-Time Learnability for Neural Networks

Neural Information Processing Systems (NeurIPS) 2017

Surbhi Goel*, Varun Kanade*, Adam R. Klivans*, Justin Thaler*

Reliably Learning ReLU in Polynomial Time

Conference on Learning Theory (COLT) 2017

WORKSHOP PAPERS

GuanWen Qiu, Da Kuang, **Surbhi Goel**

Complexity Matters: Feature Learning in the Presence of Spurious Correlations

Mathematics of Modern Machine Learning, Neural Information Processing Systems (NeurIPS) 2023

Bingbin Liu, Jordan T. Ash, **Surbhi Goel**, Akshay Krishnamurthy, Cyril Zhang

Exposing Attention Glitches with Flip-Flop Language Modeling

Challenges of Deploying Generative AI, International Conference on Machine Learning (ICML) 2023

Knowledge and Logical Reasoning in the Era of Data-driven Learning, International Conference on Machine Learning (ICML) 2023

Jessica Hoffmann, Soumya Basu, **Surbhi Goel**, Constantine Caramanis

Disentangling Mixtures of Epidemics on Graphs

Graph Representation Learning, Neural Information Processing Systems (NeurIPS) 2019

Surbhi Goel*, Adam R. Klivans*

Learning Depth-Three Neural Networks in Polynomial Time

Deep Learning: Bridging Theory and Practice, Neural Information Processing Systems (NeurIPS) 2017

Surbhi Goel*, Varun Kanade*, Adam R. Klivans*, Justin Thaler*

Reliably Learning ReLU in Polynomial Time

Optimization for Machine Learning (OPT), Neural Information Processing Systems (NeurIPS) 2016

Selected for an oral presentation

UNPUBLISHED MANUSCRIPTS

Surbhi Goel*, Rina Panigrahy*

Learning Two layer Networks with Multinomial Activation and High Thresholds

Manuscript, 2019

Matthew Jordan, Naren Manoj, **Surbhi Goel**, Alexandros Dimakis

Quantifying Perceptual Distortion of Adversarial Examples

Manuscript, 2019

Simon Du*, **Surbhi Goel***

Improved Learning of One-hidden-layer Convolutional Neural Networks with Overlaps

Manuscript, 2018.

INVITED TALKS

Understanding Training Dynamics in Deep Learning using Simplified Models

Optimization Seminar at UPenn

March 2024

How do Large Language Models Think?

AI for Executives: Daylong Immersion at Penn Engineering

May 2024

Women in Data Science at UPenn

February 2024

Beyond Worst-case Sequential Prediction: Adversarial Robustness via Abstention

Emerging Paradigms of Generalization Workshop at Simons Institute

September 2024

Rutgers/DIMACS Theory Seminar

April 2024

MINDS Seminar at JHU

March 2024

EnCORE Workshop at IPAM, UCLA

March 2024

Theory Seminar at UPenn

November 2023

<i>Alg-ML Seminar at Princeton</i>	<i>November 2023</i>
<i>BLISS Seminar at UC Berkeley</i>	<i>October 2023</i>
<i>Math Machine Learning Seminar at MPI MIS + UCLA</i>	<i>August 2023</i>
<i>FODSI Workshop on Computational Complexity of Statistical Problems at MIT</i>	<i>June 2023</i>

Thinking fast with Transformers - Algorithmic Reasoning via Shortcuts

<i>Deep Learning Down Under Workshop, Lorne, Australia</i>	<i>January 2024</i>
<i>IFML Workshop on Generative AI at UT Austin</i>	<i>November 2023</i>
<i>Youth in High Dimensions, Trieste, Italy</i>	<i>May 2023</i>
<i>MaD Seminar at NYU</i>	<i>April 2023</i>
<i>ASSET Seminar at UPenn</i>	<i>April 2023</i>

Sparse Feature Emergence in Deep Learning

<i>Symposium on New Directions in Theoretical Machine Learning [slides]</i>	<i>September 2022</i>
---	-----------------------

What Functions do Self-attention Blocks Prefer to Represent?

Demystifying Attention-based Architectures in Deep Learning

<i>Joint IFML/Data-Driven Decision Processes Workshop at Simons Institute</i>	<i>October 2022</i>
<i>ML Foundations Seminar at MSR Redmond</i>	<i>August 2022</i>
<i>Workshop on Algorithms for Learning and Economics (WALE) in Greece</i>	<i>June 2022</i>
<i>ML Symposium at USC</i>	<i>December 2021</i>
<i>ELLIS Talk Series at IST Austria</i>	<i>December 2021</i>
<i>Learning Theory Workshop at Google</i>	<i>October 2021</i>

The Hidden Progress Behind Loss Curves

<i>Workshop on Learning: Optimization and Stochastics at EPFL</i>	<i>July 2022</i>
---	------------------

Principled Algorithm Design in the Era of Deep Learning

<i>CS/CSE Colloquium at NYU Courant/Tandon</i>	<i>April 2022</i>
<i>CS Colloquium at UW-Madison</i>	<i>March 2022</i>
<i>CS Colloquium at Halicioglu Data Science Institute, UCSD</i>	<i>March 2022</i>
<i>CS Colloquium at UMD</i>	<i>February 2022</i>
<i>SCS Talk at CMU</i>	<i>February 2022</i>
<i>CS Colloquium at Duke</i>	<i>February 2022</i>
<i>CIS Colloquium at UPenn</i>	<i>February 2022</i>
<i>CS Colloquium at Cornell</i>	<i>February 2022</i>
<i>Talks at TTIC</i>	<i>February 2022</i>

Computational Barriers For Learning Some Generalized Linear Models

<i>Information-Computation Trade-offs Workshop at Simons Institute [video][slides]</i>	<i>September 2021</i>
--	-----------------------

Computational Complexity of ReLU Regression

<i>The Multifaceted Complexity of Machine Learning Workshop at IMSI [video]</i>	<i>April 2021</i>
---	-------------------

Computational Complexity of Learning Neural Networks over Gaussian Marginals

<i>Statistics Seminar at Stanford University</i>	<i>July 2021</i>
<i>SILO Seminar at UW-Madison</i>	<i>January 2021</i>
<i>TOC Colloquium at MIT</i>	<i>December 2020</i>
<i>IDEAL Seminar at TTIC</i>	<i>November 2020</i>
<i>ARC Colloquium at Georgia Tech</i>	<i>November 2020</i>
<i>ML Theory Seminar at Harvard University [video]</i>	<i>October 2020</i>
<i>Algorithms Seminar at Duke University</i>	<i>October 2020</i>
<i>MIC Seminar at NYU</i>	<i>May 2020</i>

Approximation Schemes for ReLU Regression*Deep Learning Program Reunion at Simons Institute**August 2020***Provably Efficient Algorithms for Learning Neural Networks***Microsoft Research New York**February 2020**Microsoft Research New England**February 2020**Microsoft Research Redmond**February 2020***Exploring Surrogate Losses for Learning Neural Networks***TTIC Young Researcher Seminar Series**December 2019***Efficiently Learning Simple Neural Networks***Rising Star in ML Talk at University of Maryland**September 2019***Learning Ising Models with Independent Failures***Research Fellows Talk at Simons Institute**July 2019***Efficiently Learning Simple Convolutional Networks***China Theory Week at Tsinghua University**September 2018***Learning One Convolutional Layer with Overlapping Patches***Google Research Theory Reading Group**June 2018***WORK EXPERIENCE**

Google, Mountain View CA*Research Intern**May - August 2018**Supervisor: Rina Panigrahy***Dell, Round Rock TX***Research Intern**June - August 2017***Google, New York, NY***Research Intern**May - August 2016**Supervisor: Natalia Ponomareva***Google, Mountain View CA***Software Engineering Intern**May - August 2014**Supervisor: Neha Jha***University of Michigan, Ann Arbor MI***Research Scholar**May - July 2013**Supervisor: Atul Prakash***TEACHING**

CIS 5200: Machine Learning*Co-instructor with Eric Wong**Spring 2024**University of Pennsylvania***CIS 7000: Foundations of Modern ML - Theory and Empirics***Instructor**Fall 2023**University of Pennsylvania***CIS 5200: Machine Learning***Co-instructor with Eric Wong**Spring 2023**University of Pennsylvania*

OUTREACH

Co-founder and Organizing Committee Member <i>Learning Theory Alliance (LeT-All)</i>	2020-Present
Co-organized the <i>Fall 2023 Mentoring Workshop</i>	
Co-organized the <i>Fall 2022 Mentoring Workshop</i> in collaboration with FODSI	
Co-organized the <i>COLT 2022 Mentoring Panel</i>	
Co-organized the <i>ALT 2022 Mentoring Workshop</i>	
Co-organized the <i>Graduate Applications Support Program</i> in collaboration with WiML-T	
Co-organized the <i>COLT 2021 Mentoring Workshop</i>	
Co-organized the <i>ALT 2021 Mentoring Workshop</i>	
Mentor <i>Women in Machine Learning Theory (WiML-T) Mentoring Program</i>	2021-Present
<i>UT Austin's Women in CS (GWC-WiCS) Mentoring Program</i>	2018-19
Panelist <i>New in ML Workshop, NeurIPS 2023</i>	Decemeber 2023
<i>WiML Un-Workshop, ICML 2022</i>	July 2022
<i>New Horizons in Theoretical Computer Science</i>	June 2022
<i>VMware Nirman for Women in Tech</i>	January 2021

SERVICE ROLES

Steering Committee Member <i>Association for Algorithmic Learning Theory</i>	2024-Present
Action Editor <i>Transactions on Machine Learning Research</i>	2024-Present
Co-treasurer <i>Association for Computational Learning</i>	2024-Present
Workshop Co-organizer <i>Transformers as a Computational Model, Simons Institute's Special Year on LLMs and Transformers</i>	2024-Present
Workshop Co-organizer <i>Unknown Futures of Generalization, Simons Institute's Program on Modern Paradigms in Generalization</i>	2023-Present
Program Co-organizer <i>Simons Institute's Special Year on LLMs and Transformers</i>	2023-Present
Office Hours Co-chair <i>International Conference on Learning Representations (ICLR) 2024</i>	2023-24
Workshop Reviewing Committee <i>International Conference on Machine Learning (ICML) 2024</i>	2024
Workshop Co-organizer <i>Mathematics of Modern Machine Learning (M3L) at NeurIPS 2023</i>	2023
Virtual Experience Co-chair <i>Conference on Learning Theory (COLT) 2023</i>	2023

Online Experience Co-chair	2021
<i>Conference on Learning Theory (COLT) 2021</i>	
Co-organized the virtual part of the hybrid conference, including the 2-day virtual-only program	
Seminar Co-organizer	2020-21
<i>One World Machine Learning Seminar Series</i>	
Treasurer	2016-17
<i>Graduate Representative Association of Computer Sciences (GRACS) 2024</i>	
Program Committee	
<i>International Conference on Algorithmic Learning Theory (ALT)</i>	2021/22/23
<i>Conference on Learning Theory (COLT)</i>	2021/22
<i>International Conference on Artificial Intelligence and Statistics (AISTATS) (area chair)</i>	2023
<i>Neural Information Processing Systems (NeurIPS) (area chair)</i>	2023/24
<i>International Conference on Algorithmic Learning Theory (ALT) (senior program committee)</i>	2024
<i>Conference on Learning Theory (COLT) (senior program committee)</i>	2024
Conference Reviewing	
<i>Symposium on Theory of Computing (STOC)</i>	2019/20/21
<i>Neural Information Processing Systems (NeurIPS)</i>	2018 (top 30%)/20/21
<i>Conference on Learning Theory (COLT)</i>	2018/19/20
<i>International Conference on Learning Representations (ICLR)</i>	2019/20/23
<i>Symposium on Discrete Algorithms (SODA)</i>	2020/23
<i>Foundations of Computer Science (FOCS)</i>	2020/22
<i>International Conference on Machine Learning (ICML)</i>	2019 (top 5%)
Journal Reviewing	
<i>TheoretiCS</i>	2024
<i>Journal of Machine Learning Research</i>	2021/22
<i>IEEE Transactions on Information Theory</i>	2020