SURBHI GOEL

https://www.surbhigoel.com
[first name][last initial]@cis.upenn.edu

EDUCATION

The University of Texas at Austin August 2015 - June 2020
M.S. and Ph.D. in Computer Science
Advisor: Adam R. Klivans
Dissertation: Towards Provably Efficient Algorithms for Learning Neural Networks
Committee: Alex Dimakis, Raghu Meka, Eric Price

Indian Institute of Technology, Delhi July 2011 - May 2015
B.Tech. in Computer Science and Engineering

APPOINTMENTS

University of Pennsylvania, Philadelphia, PA January 2023 *(expected)*
Magerman Term Assistant Professor, Computer and Information Science

Microsoft Research, New York, NY July 2020 - November 2022 *(expected)*
Postdoctoral Researcher, Machine Learning Group

Institute for Advanced Study, Princeton, NJ January - May 2020
Visiting Graduate Student, Theoretical Machine Learning Program

Simons Institute for Theory of Computing, Berkeley, CA May - August 2019
Research Fellow, Foundations of Deep Learning Program

RESEARCH INTERESTS

My research is on the theoretical aspects of the modern practice of machine learning, where my goal is to develop the next generation of principled machine learning methods. In the pursuit of this goal, my work focuses on quantifying the computational and statistical aspects of state-of-the-art deep learning methods, and expanding the toolbox of current algorithms using new theoretically grounded insights.

AWARDS AND FELLOWSHIPS

2020 Bert Kay Dissertation Award for best dissertation in CS at UT Austin
2019 Rising Stars in ML by University of Maryland
2019 Rising Stars in EECS by UIUC
2019 The University of Texas at Austin Graduate Dean's Prestigious Fellowship Supplement
2019 J.P. Morgan AI PhD Fellowship
2019 Simons-Berkeley Research Fellowship for Foundations of Deep Learning program
2018 The University of Texas at Austin Graduate Continuing Bruton Fellowship
2017 The University of Texas at Austin Graduate School Summer Fellowship
2015 ICIM Stay Ahead Award for Undergraduate Thesis
2015 Suresh Chandra Memorial Trust Award for Undergraduate Thesis
2011 Aditya Birla Scholarship awarded to 12 students from all over India
2011 OPJEM Scholarship awarded to 1 out of 850 students in the batch at IIT Delhi
2011 All India Rank 37 (Rank 2 in girls) in IITJEE among 450,000 students
2010 National Mathematics Olympiad finalist (1 out of 30 from all over India)

PUBLICATIONS

* *indicates α-β (alphabetical) ordering.*

THESIS

Surbhi Goel
Towards Provably Efficient Algorithms for Learning Neural Networks
The University of Texas at Austin, 2020
Received the Bert Kay dissertation award

CONFERENCE PAPERS

Benjamin L. Edelman*, **Surbhi Goel***, Sham M. Kakade*, Cyril Zhang *
Inductive Biases and Variable Creation in Self-Attention Mechanisms
International Conference on Machine Learning (ICML) 2022

Nikunj Saunshi, Jordan T. Ash, **Surbhi Goel**, Dipendra Misra, Cyril Zhang, Sanjeev Arora, Sham M. Kakade, Akshay Krishnamurthy
Understanding Contrastive Learning Requires Incorporating Inductive Biases
International Conference on Machine Learning (ICML) 2022

Jordan T. Ash, Cyril Zhang, **Surbhi Goel**, Akshay Krishnamurthy, Sham M. Kakade
Anti-Concentrated Confidence Bonuses For Scalable Exploration
International Conference on Learning Representations (ICLR) 2022

Jordan T. Ash*, **Surbhi Goel***, Akshay Krishnamurthy*, Dipendra Misra*
Investigating the Role of Negatives in Contrastive Representation Learning
International Conference on Artificial Intelligence and Statistics (AISTATS) 2022

Jordan T. Ash, **Surbhi Goel**, Akshay Krishnamurthy, Sham M. Kakade
Gone Fishing: Neural Active Learning with Fisher Embeddings
Neural Information Processing Systems (NeurIPS) 2021

Naman Agarwal*, **Surbhi Goel***, Cyril Zhang*
Acceleration via Fractal Learning Rate Schedules
International Conference on Machine Learning (ICML) 2021

Anthimos-Vardis Kandiros, Yuval Dagan, Nishanth Dikkala, **Surbhi Goel**, Constantinos Daskalakis
Statistical Estimation from Dependent Data
International Conference on Machine Learning (ICML) 2021

Surbhi Goel*, Adam R. Klivans*, Pasin Manurangsi*, Daniel Reichman*
Tight Hardness Results for Learning One-Layer ReLU Networks
Innovations in Theoretical Computer Science (ITCS) 2021

Surbhi Goel*, Adam R. Klivans*, Frederic Koehler*
From Boltzmann Machines to Neural Networks and Back Again
Neural Information Processing Systems (NeurIPS) 2020

Surbhi Goel*, Aravind Gollakota*, Adam R., Klivans*
Statistical-Query Lower Bounds via Functional Gradients

Neural Information Processing Systems (NeurIPS) 2020

Surbhi Goel*, Aravind Gollakota*, Zhihan Jin*, Sushrut Karmalkar*, Adam R. Klivans*
Superpolynomial Lower Bounds for Learning One-Layer Neural Networks using Gradient Descent
International Conference on Machine Learning (ICML) 2020

Omar Montasser, **Surbhi Goel**, Ilias Diakonikolas, Nathan Srebro
Efficiently Learning Adversarially Robust Halfspaces with Noise
International Conference on Machine Learning (ICML) 2020

Jessica Hoffmann, Soumya Basu, **Surbhi Goel**, Constantine Caramanis
Learning Mixtures of Graphs from Epidemic Cascades
International Conference on Machine Learning (ICML) 2020

Ilias Diakonikolas*, **Surbhi Goel***, Sushrut Karmalkar*, Adam R. Klivans*, Mahdi Soltanolkotabi*
Approximation Schemes for ReLU Regression
Conference on Learning Theory (COLT) 2020

Surbhi Goel
Learning Ising and Potts Models with Latent Variables
International Conference on Artificial Intelligence and Statistics (AISTATS) 2020

Surbhi Goel*, Sushrut Karmalkar*, Adam R. Klivans*
Time/Accuracy Trade-offs for Learning a ReLU with respect to Gaussian Marginals
Neural Information Processing Systems (NeurIPS) 2019
Selected for a spotlight presentation

Surbhi Goel*, Daniel Kane*, Adam R. Klivans*
Learning Ising Models with Independent Failures
Conference on Learning Theory (COLT) 2019

Surbhi Goel*, Adam R. Klivans*
Learning Neural Networks with Two Nonlinear Layers in Polynomial Time
Conference on Learning Theory (COLT) 2019

Surbhi Goel*, Adam R. Klivans*, Raghu Meka*
Learning One Convolutional Layer with Overlapping Patches
International Conference on Machine Learning (ICML) 2018
Selected for a full oral presentation

Surbhi Goel*, Adam R. Klivans*
Eigenvalue Decay Implies Polynomial-Time Learnability for Neural Networks
Neural Information Processing Systems (NeurIPS) 2017

Surbhi Goel*, Varun Kanade*, Adam R. Klivans*, Justin Thaler*
Reliably Learning ReLU in Polynomial Time
Conference on Learning Theory (COLT) 2017

WORKSHOP PAPERS

Jessica Hoffmann, Soumya Basu, **Surbhi Goel**, Constantine Caramanis

Disentangling Mixtures of Epidemics on Graphs
Graph Representation Learning, Neural Information Processing Systems (NeurIPS) 2019

Surbhi Goel*, Adam R. Klivans*
Learning Depth-Three Neural Networks in Polynomial Time
Deep Learning: Bridging Theory and Practice, Neural Information Processing Systems (NeurIPS) 2017

Surbhi Goel*, Varun Kanade*, Adam R. Klivans*, Justin Thaler*
Reliably Learning ReLU in Polynomial Time
Optimization for Machine Learning (OPT), Neural Information Processing Systems (NeurIPS) 2016
Selected for an oral presentation

WORKING PAPERS

Surbhi Goel*, Sham M. Kakade*, Adam T. Kalai*, Cyril Zhang *
Recurrent Convolutional Neural Networks Learn Succinct Learning Algorithms
In submission, 2022

Boaz Barak*, Benjamin L. Edelman*, **Surbhi Goel***, Sham M. Kakade*, Eran Malach*, Cyril Zhang*
Hidden Progress in Deep Learning: SGD Learns Parities Near the Computational Limit
In submission, 2022

MANUSCRIPTS

Surbhi Goel*, Rina Panigrahy*
Learning Two layer Networks with Multinomial Activation and High Thresholds
Manuscript, 2019

Matthew Jordan, Naren Manoj, **Surbhi Goel**, Alexandros Dimakis
Quantifying Perceptual Distortion of Adversarial Examples
Manuscript, 2019

Simon Du*, **Surbhi Goel***
Improved Learning of One-hidden-layer Convolutional Neural Networks with Overlaps
Manuscript, 2018.

INVITED TALKS

Sparse Feature Emergence in Deep Learning
Simons Foundation Symposium on New Directions in Theoretical Machine Learning *September 2022*

What Functions do Self-attention Blocks Prefer to Represent?
Demystifying Attention-based Architectures in Deep Learning
ML Foundations Seminar at MSR Redmond *August 2022*
Workshop on Algorithms for Learning and Economics (WALE) in Greece *June 2022*
ML Symposium at USC *December 2021*
ELLIS Talk Series at IST Austria *December 2021*
Learning Theory Workshop at Google *October 2021*

The Hidden Progress Behind Loss Curves
Workshop on Learning: Optimization and Stochastics at EPFL *July 2022*

Pricipled Algorithm Design in the Era of Deep Learning
CS/CSE Colloquium at NYU Courant/Tandon *April 2022*
CS Colloquium at UW-Madison *March 2022*

CS Colloquium at Halicioglu Data Science Institute, UCSD — *March 2022*
CS Colloquium at UMD — *February 2022*
SCS Talk at CMU — *February 2022*
CS Colloquium at Duke — *February 2022*
CIS Colloquium at UPenn — *February 2022*
CS Colloquium at Cornell University — *February 2022*
Talks at TTIC — *February 2022*

Computational Barriers For Learning Some Generalized Linear Models
Information-Computation Trade-offs Workshop at Simons Institute [video][slides] — *September 2021*

Computational Complexity of ReLU Regression
The Multifaceted Complexity of Machine Learning Workshop at IMSI [video] — *April 2021*

Computational Complexity of Learning Neural Networks over Gaussian Marginals
MIC Seminar at NYU — *May 2020*
Algorithms Seminar at Duke University — *October 2020*
ML Theory Seminar at Harvard University [video] — *October 2020*
ARC Colloquium at Georgia Tech — *November 2020*
IDEAL Seminar at TTIC — *November 2020*
TOC Colloquium at MIT — *December 2020*
SILO Seminar at UW-Madison — *January 2020*
Statistics Seminar at Stanford University — *July 2021*

Approximation Schemes for ReLU Regression
Deep Learning Program Reunion at Simons Institute — *August 2020*

Provably Efficient Algorithms for Learning Neural Networks
Microsoft Research New York — *February 2020*
Microsoft Research New England — *February 2020*
Microsoft Research Redmond — *February 2020*

Time/Accuracy Tradeoffs for Learning a ReLU wrt Gaussian Marginals
Spotlight Talk at Neural Information Processing Systems (NeurIPS) — *December 2019*

Exploring Surrogate Losses for Learning Neural Networks
TTIC Young Researcher Seminar Series — *December 2019*

Efficiently Learning Simple Neural Networks
Rising Star in ML Talk at University of Maryland — *September 2019*

Learning Ising Models with Independent Failures
Research Fellows Talk at Simons Institute — *July 2019*

Efficiently Learning Simple Convolutional Networks
China Theory Week at Tsinghua University — *September 2019*

Learning One Convolutional Layer with Overlapping Patches
Google Research Theory Reading Group — *June 2018*

Reliably Learning the ReLU in Polynomial Time
OPT-ML Workshop at Neural Information Processing Systems (NeurIPS) — *December 2016*

WORK EXPERIENCE

Google, Mountain View CA — May - August 2018
Research Intern — *Supervisor: Rina Panigrahy*

Dell, Round Rock TX — June - August 2017
Research Intern

Google, New York, NY — May - August 2016
Research Intern — *Supervisor: Natalia Ponomareva*

Google, Mountain View CA — May - August 2014
Software Engineering Intern — *Supervisor: Neha Jha*

University of Michigan, Ann Arbor MI — May - July 2013
Research Scholar — *Supervisor: Atul Prakash*

TEACHING EXPERIENCE

University of Texas at Austin — Spring 2018
Course: Data Mining (Hons.) — *Guest Lecture*

University of Texas at Austin — Spring 2016
Course: Distributed Computing (Hons.) — *Teaching Assistant*

University of Texas at Austin — Fall 2015
Course: Data Structures — *Teaching Assistant*

Indian Institute of Technology Delhi — Spring 2015
Course: Data Structures — *Teaching Assistant*

OUTREACH

Co-founder — 2020-Present
Learning Theory Alliance (LeT-All)

Co-organizing the Fall 2022 Mentoring Workshop
Co-organized the COLT 2022 Mentoring Panel
Co-organized the ALT 2022 Mentoring Workshop
Co-organized the Graduate Applications Support Program
Co-organized the COLT 2021 Mentoring Workshop
Co-organized the ALT 2021 Mentoring Workshop

Mentor
Women in Machine Learning Theory (WiML-T) Mentoring Program — *2021-Present*
UT Austin's Women in CS (GWC-WiCS) Mentoring Program — *2018-19*

Panelist
WiML Un-Workshop, ICML 2022 — *July 2022*
New Horizons in Theoretical Computer Science — *June 2022*
VMware Nirman for Women in Tech — *January 2021*

SERVICE ROLES

Virtual Experience Chair 2021
Conference on Learning Theory (COLT)

Co-organized the virtual part of the hybrid conference, including the 2-day virtual-only program

Co-organizer 2020-2021
One World Machine Learning Seminar Series

Treasurer 2016-17
Graduate Representative Association of Computer Sciences (GRACS)

Program Committee

International Conference on Algorithmic Learning Theory (ALT)	*2021/22*
Conference on Learning Theory (COLT)	*2021/22*
International Conference on Artificial Intelligence and Statistics (AISTATS) (area chair)	*2023*

Conference Reviewing

Symposium on Theory of Computing (STOC)	*2019/20/21*
Neural Information Processing Systems (NeurIPS)	*2018 (top 30%)/20/21*
Conference on Learning Theory (COLT)	*2018/19/20*
International Conference on Learning Representations (ICLR)	*2019/20/23*
Symposium on Discrete Algorithms (SODA)	*2020/23*
Foundations of Computer Science (FOCS)	*2020/22*
International Conference on Machine Learning (ICML)	*2019 (top 5%)*

Journal Reviewing

Journal of Machine Learning Research	*2021/22*
IEEE Transactions on Information Theory	*2020*